

✓ Importing the libraries

```
import tensorflow as tf
print(tf.__version__)
```

↗ 2.18.0

```
import cv2
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from google.colab.patches import cv2_imshow
import zipfile
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Conv2D, MaxPooling2D, Flatten, BatchNormalization
```

✓ Loading the images

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Mounted at /content/drive

```
path = '/content/drive/MyDrive/Computer Vision Masterclass/Datasets/fer_images.zip'
zip_object = zipfile.ZipFile(file=path, mode='r')
zip_object.extractall('.')
zip_object.close()
```

```
tf.keras.preprocessing.image.load_img('/content/fer2013/train/Angry/1003.jpg')
```



```
image = tf.keras.preprocessing.image.load_img('/content/fer2013/train/Happy/1.jpg')
image
```



✓ Train and test set

```
training_generator = ImageDataGenerator(rescale=1./255,
                                         rotation_range=7,
                                         horizontal_flip=True,
                                         zoom_range=0.2)
train_dataset = training_generator.flow_from_directory('/content/fer2013/train',
                                                       target_size = (48, 48),
                                                       batch_size = 16,
                                                       class_mode = 'categorical',
                                                       shuffle = True)
```

↗ Found 28709 images belonging to 7 classes.

```
train_dataset.classes
```

↗ array([0, 0, 0, ..., 6, 6, 6], dtype=int32)


```
np.unique(train_dataset.classes, return_counts=True)
```

↗ (array([0, 1, 2, 3, 4, 5, 6], dtype=int32),
array([3995, 436, 4097, 7215, 4965, 4830, 3171]))

```
train_dataset.class_indices
```

```
{'Angry': 0,
 'Disgust': 1,
 'Fear': 2,
 'Happy': 3,
 'Neutral': 4,
 'Sad': 5,
 'Surprise': 6}
```

```
sns.countplot(x = train_dataset.classes);
```



```
test_generator = ImageDataGenerator(rescale=1./255)
test_dataset = test_generator.flow_from_directory('/content/fer2013/validation',
                                                  target_size = (48, 48),
                                                  batch_size = 1,
                                                  class_mode = 'categorical',
                                                  shuffle = False)
```

```
Found 3589 images belonging to 7 classes.
```

✓ Building and training the convolutional neural network

```
2*2*2*32
```

```
256
```

```
num_detectors = 32
num_classes = 7
width, height = 48, 48
epochs = 70
```

```
network = Sequential() # i gonna define a sequence of layers
```

```
network.add(Conv2D(num_detectors, (3,3), activation='relu', padding = 'same', input_shape = (width, height, 3)))
network.add(BatchNormalization())
network.add(Conv2D(num_detectors, (3,3), activation='relu', padding = 'same'))
network.add(BatchNormalization())
network.add(MaxPooling2D(pool_size=(2,2)))
network.add(Dropout(0.2))
```

```
network.add(Conv2D(2*num_detectors, (3,3), activation='relu', padding = 'same'))
network.add(BatchNormalization())
network.add(Conv2D(2*num_detectors, (3,3), activation='relu', padding = 'same'))
network.add(BatchNormalization())
network.add(MaxPooling2D(pool_size=(2,2)))
network.add(Dropout(0.2))
```

```
network.add(Conv2D(2*2*num_detectors, (3,3), activation='relu', padding = 'same'))
network.add(BatchNormalization())
network.add(Conv2D(2*2*num_detectors, (3,3), activation='relu', padding = 'same'))
network.add(BatchNormalization())
network.add(MaxPooling2D(pool_size=(2,2)))
network.add(Dropout(0.2))
```

```
network.add(Conv2D(2*2*2*num_detectors, (3,3), activation='relu', padding = 'same'))
network.add(BatchNormalization())
network.add(Conv2D(2*2*2*num_detectors, (3,3), activation='relu', padding = 'same'))
```

```
network.add(Conv2D(2 * num_detectors, (3,3), activation='relu', padding='same'))
network.add(BatchNormalization())
network.add(MaxPooling2D(pool_size=(2,2)))
network.add(Dropout(0.2))

network.add(Flatten())

network.add(Dense(2 * num_detectors, activation='relu'))
network.add(BatchNormalization())
network.add(Dropout(0.2))

network.add(Dense(2 * num_detectors, activation='relu'))
network.add(BatchNormalization())
network.add(Dropout(0.2))

network.add(Dense(num_classes, activation='softmax'))
print(network.summary())
```

```

    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
    Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 48, 48, 32)	896
batch_normalization (BatchNormalization)	(None, 48, 48, 32)	128
conv2d_1 (Conv2D)	(None, 48, 48, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None, 48, 48, 32)	128
max_pooling2d (MaxPooling2D)	(None, 24, 24, 32)	0
dropout (Dropout)	(None, 24, 24, 32)	0
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None, 24, 24, 64)	256
conv2d_3 (Conv2D)	(None, 24, 24, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None, 24, 24, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
conv2d_4 (Conv2D)	(None, 12, 12, 128)	73,856
batch_normalization_4 (BatchNormalization)	(None, 12, 12, 128)	512
conv2d_5 (Conv2D)	(None, 12, 12, 128)	147,584
batch_normalization_5 (BatchNormalization)	(None, 12, 12, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
dropout_2 (Dropout)	(None, 6, 6, 128)	0
conv2d_6 (Conv2D)	(None, 6, 6, 256)	295,168
batch_normalization_6 (BatchNormalization)	(None, 6, 6, 256)	1,024
conv2d_7 (Conv2D)	(None, 6, 6, 256)	590,080
batch_normalization_7 (BatchNormalization)	(None, 6, 6, 256)	1,024
max_pooling2d_3 (MaxPooling2D)	(None, 3, 3, 256)	0
dropout_3 (Dropout)	(None, 3, 3, 256)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 64)	147,520
batch_normalization_8 (BatchNormalization)	(None, 64)	256
dropout_4 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 64)	4,160
batch_normalization_9 (BatchNormalization)	(None, 64)	256
dropout_5 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 7)	455

Total params: 1,328,743 (5.07 MB)
 Trainable params: 1,326,567 (5.06 MB)
 Non-trainable params: 2,176 (8.50 KB)
 None

```

network.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

```

network.fit(train_dataset, epochs=epochs)

```

```

1795/1795 ————— 45s 25ms/step - accuracy: 0.6994 - loss: 0.8259
Epoch 39/70
1795/1795 ————— 44s 24ms/step - accuracy: 0.7000 - loss: 0.8263
Epoch 40/70
1795/1795 ————— 81s 24ms/step - accuracy: 0.7054 - loss: 0.8183
Epoch 41/70
1795/1795 ————— 42s 24ms/step - accuracy: 0.6944 - loss: 0.8335
Epoch 42/70
1795/1795 ————— 82s 24ms/step - accuracy: 0.7022 - loss: 0.8177
Epoch 43/70
1795/1795 ————— 45s 25ms/step - accuracy: 0.7064 - loss: 0.8067
Epoch 44/70
1795/1795 ————— 45s 25ms/step - accuracy: 0.7161 - loss: 0.7840
Epoch 45/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7147 - loss: 0.7920
Epoch 46/70
1795/1795 ————— 83s 24ms/step - accuracy: 0.7182 - loss: 0.7696
Epoch 47/70
1795/1795 ————— 44s 24ms/step - accuracy: 0.7184 - loss: 0.7840
Epoch 48/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7013 - loss: 0.8096
Epoch 49/70
1795/1795 ————— 82s 24ms/step - accuracy: 0.7169 - loss: 0.7778
Epoch 50/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7166 - loss: 0.7842
Epoch 51/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7178 - loss: 0.7736
Epoch 52/70
1795/1795 ————— 44s 24ms/step - accuracy: 0.7313 - loss: 0.7476
Epoch 53/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7231 - loss: 0.7654
Epoch 54/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7249 - loss: 0.7600
Epoch 55/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7287 - loss: 0.7474
Epoch 56/70
1795/1795 ————— 82s 24ms/step - accuracy: 0.7259 - loss: 0.7492
Epoch 57/70
1795/1795 ————— 42s 24ms/step - accuracy: 0.7310 - loss: 0.7455
Epoch 58/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7254 - loss: 0.7504
Epoch 59/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7328 - loss: 0.7368
Epoch 60/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7360 - loss: 0.7258
Epoch 61/70
1795/1795 ————— 83s 25ms/step - accuracy: 0.7381 - loss: 0.7175
Epoch 62/70
1795/1795 ————— 42s 24ms/step - accuracy: 0.7404 - loss: 0.7111
Epoch 63/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7421 - loss: 0.7054
Epoch 64/70
1795/1795 ————— 43s 24ms/step - accuracy: 0.7467 - loss: 0.7088
Epoch 65/70
1795/1795 ————— 45s 25ms/step - accuracy: 0.7500 - loss: 0.6920
Epoch 66/70
1795/1795 ————— 81s 24ms/step - accuracy: 0.7515 - loss: 0.6956
Epoch 67/70

```

Saving and loading the model

```

model_json = network.to_json()
with open('network_emotions.json', 'w') as json_file:
    json_file.write(model_json)

```

```

from keras.models import save_model
network_saved = save_model(network, '/content/weights_emotions.hdf5')

```

Considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(m`

```

from tensorflow.keras.models import model_from_json, Sequential

```

```

# Load the model architecture from the JSON file
with open('/content/network_emotions.json', 'r') as json_file:
    json_saved_model = json_file.read()

```

```

network_loaded = model_from_json(
    json_saved_model,
    custom_objects={'Sequential': Sequential}
)

```

```

# Load the model weights from the HDF5 file

```

```
network_loaded.load_weights('/content/drive/MyDrive/Computer Vision Masterclass/Weights/weights_emotions.hdf5')
```

```
network_loaded.compile(loss='categorical_crossentropy', optimizer='Adam', metrics=['accuracy'])
```

```
network_loaded.summary()
```



Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 48, 48, 32)	896
batch_normalization (BatchNormalization)	(None, 48, 48, 32)	128
conv2d_1 (Conv2D)	(None, 48, 48, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None, 48, 48, 32)	128
max_pooling2d (MaxPooling2D)	(None, 24, 24, 32)	0
dropout (Dropout)	(None, 24, 24, 32)	0
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None, 24, 24, 64)	256
conv2d_3 (Conv2D)	(None, 24, 24, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None, 24, 24, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
conv2d_4 (Conv2D)	(None, 12, 12, 128)	73,856
batch_normalization_4 (BatchNormalization)	(None, 12, 12, 128)	512
conv2d_5 (Conv2D)	(None, 12, 12, 128)	147,584
batch_normalization_5 (BatchNormalization)	(None, 12, 12, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
dropout_2 (Dropout)	(None, 6, 6, 128)	0
conv2d_6 (Conv2D)	(None, 6, 6, 256)	295,168
batch_normalization_6 (BatchNormalization)	(None, 6, 6, 256)	1,024
conv2d_7 (Conv2D)	(None, 6, 6, 256)	590,080
batch_normalization_7 (BatchNormalization)	(None, 6, 6, 256)	1,024
max_pooling2d_3 (MaxPooling2D)	(None, 3, 3, 256)	0
dropout_3 (Dropout)	(None, 3, 3, 256)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 64)	147,520
batch_normalization_8 (BatchNormalization)	(None, 64)	256
dropout_4 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 64)	4,160
batch_normalization_9 (BatchNormalization)	(None, 64)	256
dropout_5 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 7)	455

Total params: 1,328,743 (5.07 MB)

Trainable params: 1,326,567 (5.06 MB)

✦ Evaluating the neural network

```
network_loaded.evaluate(test_dataset)
```

```
→ /usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` c]
self._warn_if_super_not_called()
3589/3589 ————— 17s 4ms/step - accuracy: 0.5804 - loss: 1.3395
[1.5148499011993408, 0.5778768658638]
```

```
predictions = network_loaded.predict(test_dataset)
predictions
```

```
→ 3589/3589 ————— 12s 3ms/step
array([[9.1399026e-01, 5.2431659e-03, 2.0064995e-02, ..., 7.5534321e-03,
        3.9156068e-02, 1.2308669e-02],
       [8.3489537e-01, 1.3114783e-04, 1.6049375e-01, ..., 8.8334654e-04,
        3.4225823e-03, 1.4467943e-07],
       [9.6874458e-01, 7.4295519e-04, 1.7095286e-02, ..., 1.2566427e-03,
        7.5712553e-03, 4.9269065e-04],
       ...,
       [1.7553559e-03, 4.7019232e-05, 4.1051388e-02, ..., 1.6656774e-03,
        2.9383365e-03, 9.5245790e-01],
       [2.9105307e-03, 4.3033357e-05, 6.9266900e-02, ..., 3.6131206e-01,
        3.1379773e-03, 5.3026170e-01],
       [2.6819391e-02, 2.2662596e-03, 1.1300567e-01, ..., 4.7751591e-03,
        7.1213404e-03, 8.4558338e-01]], dtype=float32)
```

```
predictions = np.argmax(predictions, axis = 1)
predictions
```

```
→ array([0, 0, 0, ..., 6, 6, 6])
```

```
test_dataset.classes
```

```
→ array([0, 0, 0, ..., 6, 6, 6], dtype=int32)
```

```
from sklearn.metrics import accuracy_score
accuracy_score(test_dataset.classes, predictions)
```

```
→ 0.5778768459180831
```


```
test_dataset.class_indices
```

```
→ {'Angry': 0,
   'Disgust': 1,
   'Fear': 2,
   'Happy': 3,
   'Neutral': 4,
   'Sad': 5,
   'Surprise': 6}
```

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(test_dataset.classes, predictions)
cm
```

```
→ array([[276,  8, 53, 14, 64, 71,  5],
       [ 13, 36,  2,  1,  2,  1,  0],
       [ 51,  3, 244, 19, 78, 91, 42],
       [ 10,  2, 16, 775, 46, 18, 12],
       [ 87, 11, 84, 160, 115, 111, 58],
       [ 44,  1, 53, 25, 146, 319,  6],
       [  4,  3, 56, 25, 14,  5, 309]])
```

```
sns.heatmap(cm, annot=True);
```



```
from sklearn.metrics import classification_report
print(classification_report(test_dataset.classes, predictions))
```

	precision	recall	f1-score	support
0	0.57	0.56	0.57	491
1	0.56	0.65	0.61	55
2	0.48	0.46	0.47	528
3	0.76	0.88	0.82	879
4	0.25	0.18	0.21	626
5	0.52	0.54	0.53	594
6	0.72	0.74	0.73	416
accuracy			0.58	3589
macro avg	0.55	0.57	0.56	3589
weighted avg	0.56	0.58	0.56	3589

✓ Classifying one single image

```
image = cv2.imread('/content/IMG_20250506_214218.png')
cv2_imshow(image)
```



```
image.shape
```

```
(236, 225, 3)
```

```
face_detector = cv2.CascadeClassifier('/content/drive/MyDrive/Computer Vision Masterclass/Cascades/haarcascade_frontalface_default.xml')
```

```
original_image = image.copy()
faces = face_detector.detectMultiScale(original_image)
```

```
faces
```

```
array([[ 45,  38, 152, 152]], dtype=int32)
```

```
roi = image[10:200, 40:200]
cv2_imshow(roi)
```




```
roi.shape
```



```
(190, 160, 3)
```

```
roi = cv2.resize(roi, (48, 48))
cv2_imshow(roi)
```



```
roi.shape
```



```
(48, 48, 3)
```

```
roi
```



```
ndarray (48, 48, 3) show data
```



```
roi = roi / 255
roi
```



```
array([[0.9372549, 0.94509804, 0.9254902 ],
       [0.92941176, 0.94509804, 0.9254902 ],
       [0.9254902, 0.94509804, 0.91764706],
       ...,
       [0.92156863, 0.9254902, 0.90588235],
       [0.92156863, 0.9254902, 0.90588235],
       [0.92156863, 0.9254902, 0.90588235]],

       [[0.93333333, 0.94901961, 0.9254902 ],
       [0.92941176, 0.94509804, 0.92156863],
       [0.92156863, 0.9372549, 0.91372549],
       ...,
       [0.92156863, 0.9254902, 0.90588235],
       [0.92941176, 0.93333333, 0.91372549],
       [0.92156863, 0.9254902, 0.90588235]],

       [[0.94117647, 0.95686275, 0.92941176],
       [0.93333333, 0.94901961, 0.9254902 ],
       [0.9254902, 0.94117647, 0.91764706],
       ...,
       [0.91764706, 0.92156863, 0.90196078],
       [0.92941176, 0.92941176, 0.90980392],
       [0.92941176, 0.92941176, 0.90980392]],

       ...,

       [[0.80784314, 0.83137255, 0.82745098],
       [0.8, 0.82352941, 0.81960784],
       [0.78823529, 0.81568627, 0.80784314],
       ...,
       [0.81960784, 0.81960784, 0.79215686],
       [0.82352941, 0.82352941, 0.79607843],
       [0.82745098, 0.82745098, 0.8 ]],

       [[0.76862745, 0.79607843, 0.78823529],
       [0.69803922, 0.74117647, 0.70980392],
       [0.52156863, 0.6, 0.5372549 ],
       ...,
       [0.82745098, 0.81960784, 0.78431373],
       [0.82745098, 0.81568627, 0.79215686],
       [0.83529412, 0.82745098, 0.79607843]],

       [[0.55686275, 0.63529412, 0.57647059],
       [0.53333333, 0.61960784, 0.54901961],
       [0.54901961, 0.63921569, 0.56078431],
       ...,
       [0.81960784, 0.81568627, 0.77647059],
```

```
[0.82745098, 0.82352941, 0.78431373],  
[0.83137255, 0.82352941, 0.78431373]]])
```

```
roi.shape
```

```
↗ (48, 48, 3)
```

```
roi = np.expand_dims(roi, axis = 0)  
roi.shape
```

```
↗ (1, 48, 48, 3)
```

```
probs = network_loaded.predict(roi)  
probs
```

```
↗ 1/1 ————— 1s 954ms/step  
array([[1.2675990e-04, 1.9988229e-06, 5.5030128e-04, 8.9367753e-01,  
1.0431930e-01, 1.1285383e-03, 1.9558656e-04]], dtype=float32)
```

```
result = np.argmax(probs)  
result
```

```
↗ np.int64(3)
```

```
test_dataset.class_indices
```

```
↗ {'Angry': 0,  
  'Disgust': 1,  
  'Fear': 2,  
  'Happy': 3,  
  'Neutral': 4,  
  'Sad': 5,  
  'Surprise': 6}
```

classifying multiple images

```
image = cv2.imread('/content/drive/MyDrive/Computer Vision Masterclass/Images/faces_emotions.png')  
cv2_imshow(image)
```



```
faces = face_detector.detectMultiScale(image)  
faces
```

```
↗ array([[625, 49, 91, 91],  
 [224, 35, 90, 90],  
 [ 23, 41, 92, 92],  
 [420, 43, 97, 97],  
 [420, 242, 97, 97],  
 [ 18, 243, 98, 98],  
 [229, 242, 85, 85],  
 [627, 241, 91, 91]], dtype=int32)
```

```
test_dataset.class_indices.keys()
```

```
↗ dict_keys(['Angry', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad', 'Surprise'])
```

```
emotions = ['Angry', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad', 'Surprise']
```

```
for (x, y, w, h) in faces:
    cv2.rectangle(image, (x, y), (x + w, y + h), (0, 255, 0), 1)
    roi = image[y:y + h, x:x + w]
    #cv2.imshow(roi)
    roi = cv2.resize(roi, (48, 48))
    #cv2.imshow(roi)
    roi = roi / 255
    roi = np.expand_dims(roi, axis = 0)
    #print(roi.shape)
    prediction = network_loaded.predict(roi)
    #print(prediction)
    cv2.putText(image, emotions[np.argmax(prediction)], (x, y - 10), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0,255,0), 2, cv2.LINE_AA)
cv2.imshow(image)
```

```
1/1 ————— 0s 31ms/step
1/1 ————— 0s 30ms/step
1/1 ————— 0s 30ms/step
1/1 ————— 0s 31ms/step
1/1 ————— 0s 29ms/step
1/1 ————— 0s 28ms/step
1/1 ————— 0s 28ms/step
1/1 ————— 0s 44ms/step
```



✓ classifying video emotions

```
cap = cv2.VideoCapture('/content/drive/MyDrive/Computer Vision Masterclass/Videos/emotion_test01.mp4')
connected, video = cap.read()
print(connected, video.shape)
```

```
True (360, 640, 3)
```

```
# fourcc.org
save_path = '/content/drive/MyDrive/Computer Vision Masterclass/Videos/emotion_test01_result.avi'
fourcc = cv2.VideoWriter_fourcc(*'XVID')
fps = 24
output_video = cv2.VideoWriter(save_path, fourcc, fps, (video.shape[1], video.shape[0]))
```

```
while (cv2.waitKey(1) < 0):
    connected, frame = cap.read()
    if not connected:
        break
    faces = face_detector.detectMultiScale(frame, scaleFactor=1.2, minNeighbors=5, minSize=(30,30))
    if len(faces) > 0:
        for (x, y, w, h) in faces:
            frame = cv2.rectangle(frame, (x, y), (x + w, y + h), (0, 255, 0), 2)
            roi = frame[y:y + h, x:x + w]
            roi = cv2.resize(roi, (48, 48))
            roi = roi / 255
            roi = np.expand_dims(roi, axis = 0)
            prediction = network_loaded.predict(roi)

            if prediction is not None:
                result = np.argmax(prediction)
                cv2.putText(frame, emotions[result], (x, y-10), cv2.FONT_HERSHEY_SIMPLEX, 0.5, (255,255,255), 1, cv2.LINE_AA)
```